

Advanced Machine Learning

Halim BENHABILES, Ph.D

Associate professor

DSLS Team, H&E Dept., JUNIA ISEN Lille

CNRS UMR-8520 IEMN

Tél. : 33 (0)3 20 30 40 27

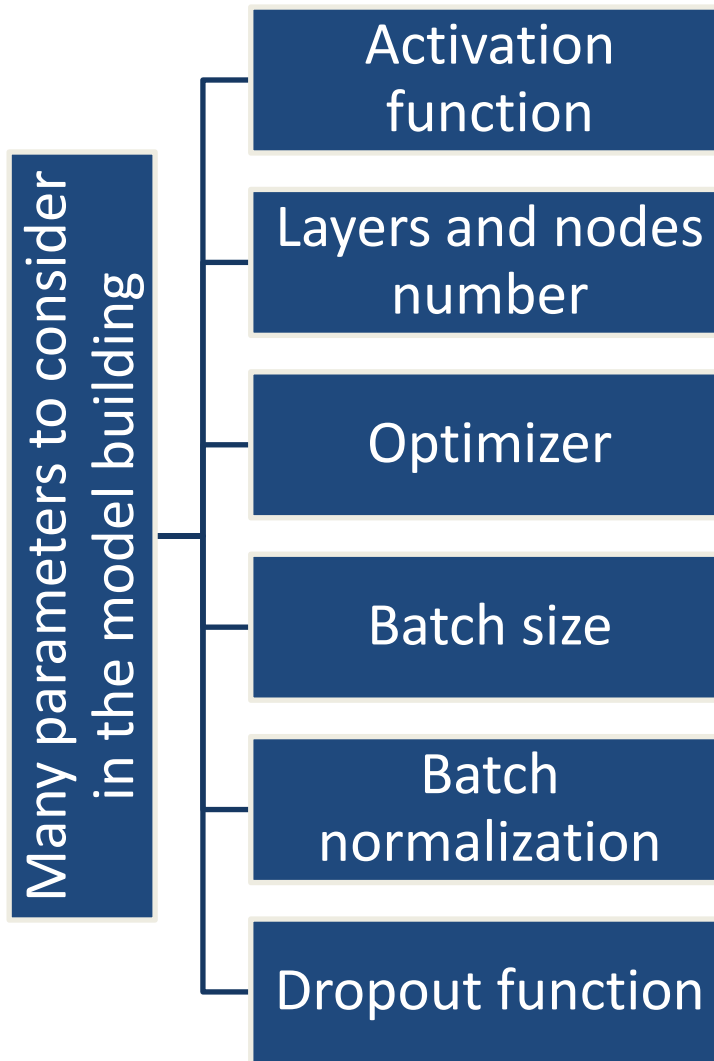
halim.benhables@junia.com

halim.benhables@iemn.fr




Institut d'Electronique, de Microélectronique et de Nanotechnologie
UMR CNRS 8520

Improving the MLP with Keras – TensorFlow core



Improving the MLP with Keras – TensorFlow core

- Activation function $\sigma(x)$ for hidden layers
 - Linear (identity) $\rightarrow x$
 - Exponential $\rightarrow e^x$
 - Elu (exponential linear unit) $\rightarrow \begin{cases} x & \text{if } x > 0 \\ \alpha(e^x - 1) & \text{else} \end{cases}$
 - Selu (scaled Elu) $\rightarrow scale * Elu(x)$
 - Relu (rectified linear unit) $\rightarrow \max(x, 0)$
 - Sigmoid $\rightarrow \frac{1}{1 + e^{-x}}$
 - Hard_sigmoid $\rightarrow \begin{cases} 0 & \text{if } x < -2.5 \\ 1 & \text{if } x > 2.5 \\ 0.2x + 0.5 & \text{if } -2.5 \leq x \leq 2.5 \end{cases}$
 - Tanh $\rightarrow \frac{e^x - e^{-x}}{e^x + e^{-x}}$
 - Softplus $\rightarrow \log(e^x + 1)$
 - Softsign $\rightarrow \frac{x}{|x| + 1}$




Improving the MLP with Keras – TensorFlow core

- Activation function $\sigma(x)$ for hidden layers

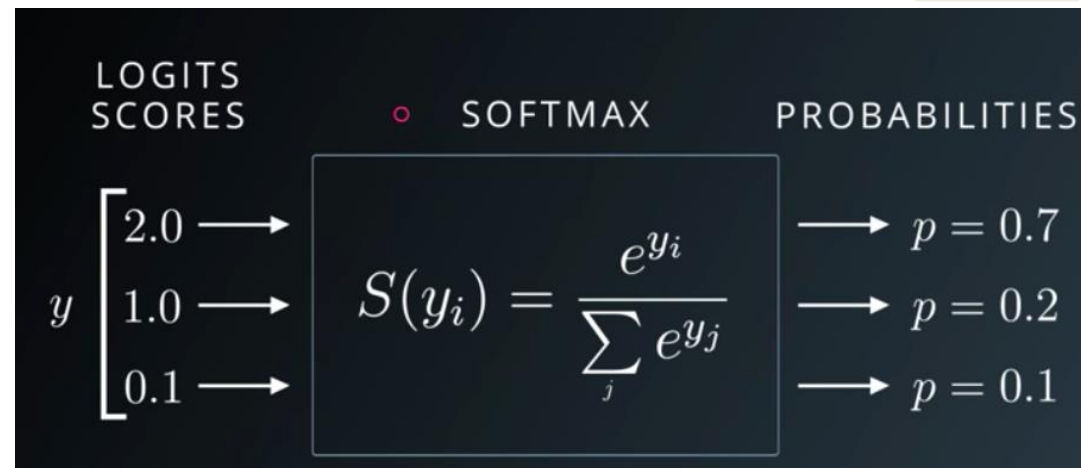


- Linear (identity) $\rightarrow x$
- Exponential $\rightarrow e^x$
- Elu (exponential linear unit) $\rightarrow \begin{cases} x & \text{if } x > 0 \\ \alpha(e^x - 1) & \text{else} \end{cases}$
- Selu (scaled Elu) $\rightarrow scale * Elu(x)$
- Relu (rectified linear unit) $\rightarrow \max(x, 0)$
- Sigmoid $\rightarrow \frac{1}{1 + e^{-x}}$
- Hard_sigmoid $\rightarrow \begin{cases} 0 & \text{if } x < -2.5 \\ 1 & \text{if } x > 2.5 \\ 0.2x + 0.5 & \text{if } -2.5 \leq x \leq 2.5 \end{cases}$
- Tanh $\rightarrow \frac{e^x - e^{-x}}{e^x + e^{-x}}$
- Softplus $\rightarrow \log(e^x + 1)$
- Softsign $\rightarrow \frac{x}{|x| + 1}$

Most used \rightarrow relu, elu, sigmoid



Improving the MLP with Keras – TensorFlow core

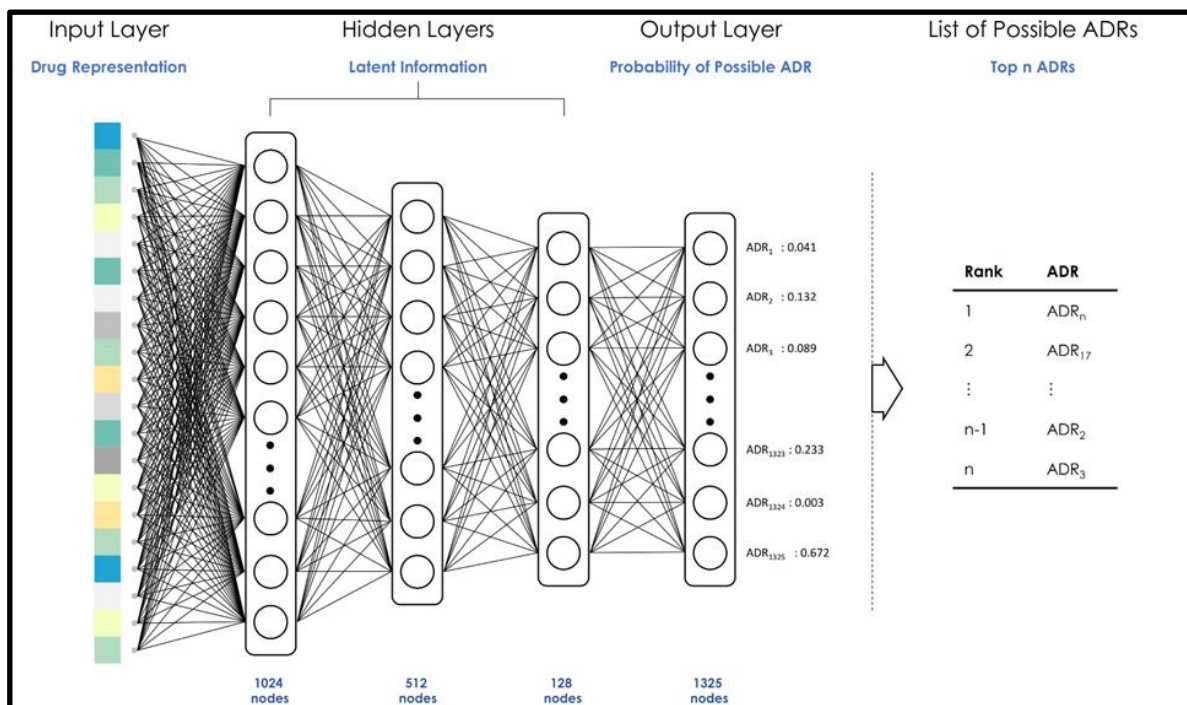

- Softmax function for **output layer** → **Transform output values into probabilities with sum=1 over the layer**



<https://medium.com/data-science-bootcamp/understand-the-softmax-function-in-minutes-f3a59641e86d>

Improving the MLP with Keras – TensorFlow core

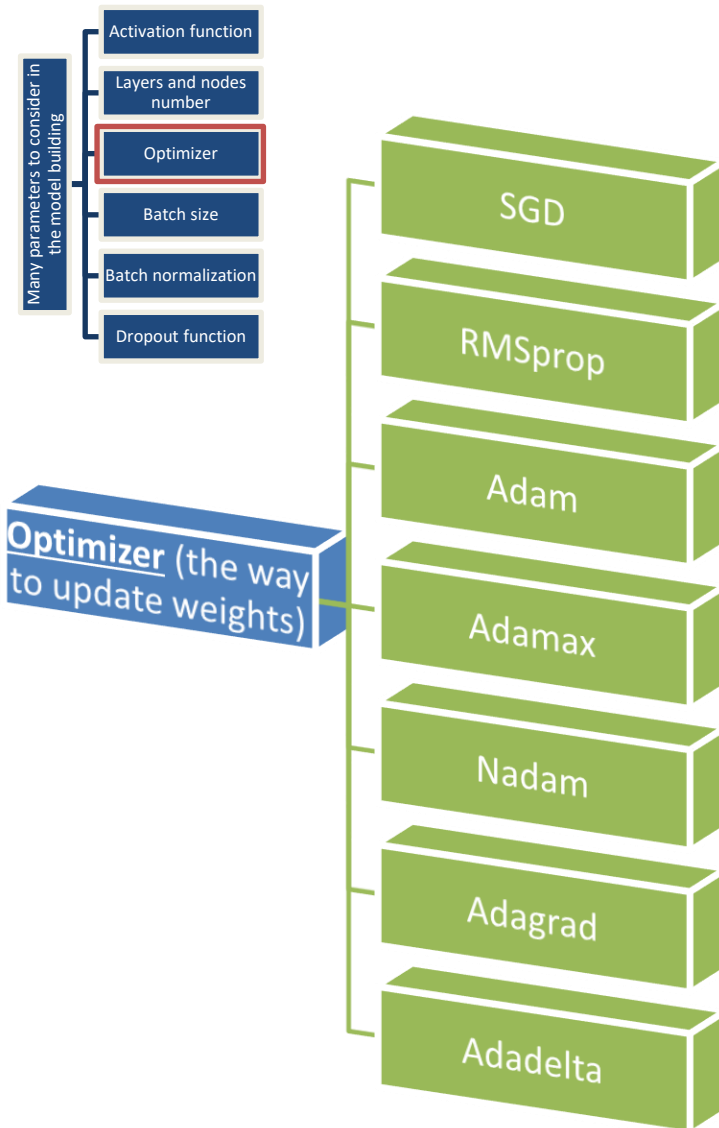
- Setting number of layers and nodes / layer (no exact answer – problem dependent)
 - Input layer → identical to the shape of your data
 - Output layer → number of classes to fix for **classification problem**
 - Hidden layers
 - Large number of layers for non linear data (time and memory limitation)
 - The first hidden layer smaller than the input layer
 - Decreasing size of layers until to reach the output (very often power of 2)



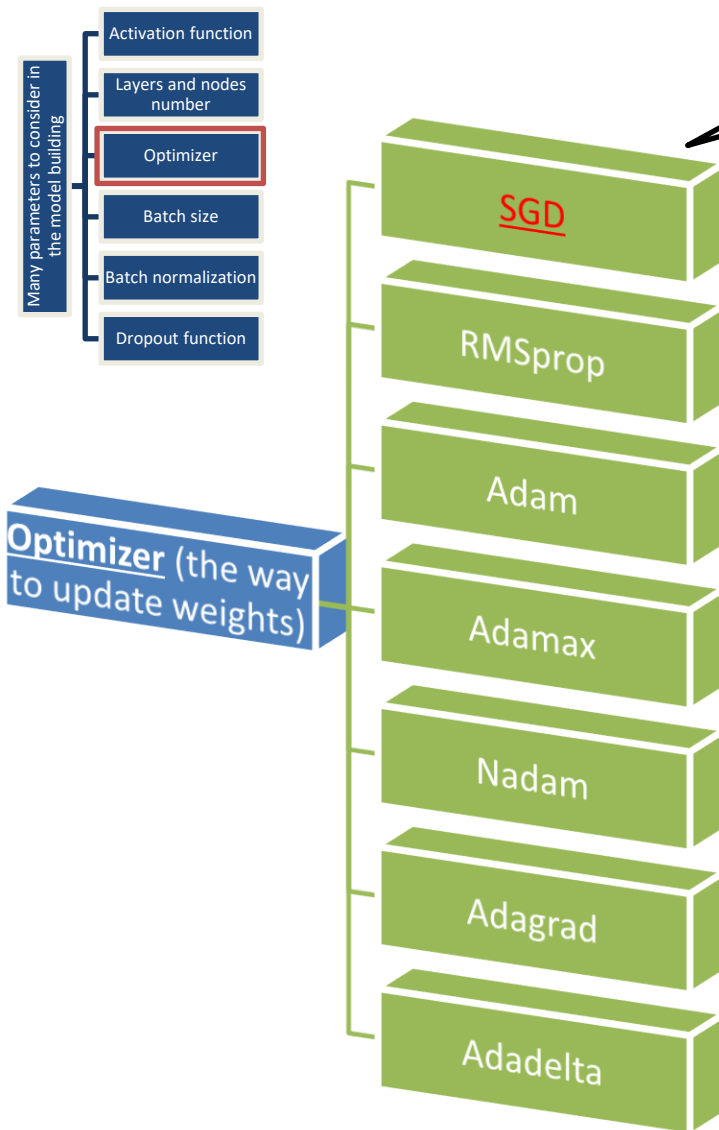
Detecting Potential Adverse Drug Reactions Using a Deep Neural Network Model -

<https://www.jmir.org/2019/2/e11016/pdf>

Improving the MLP with Keras – TensorFlow core



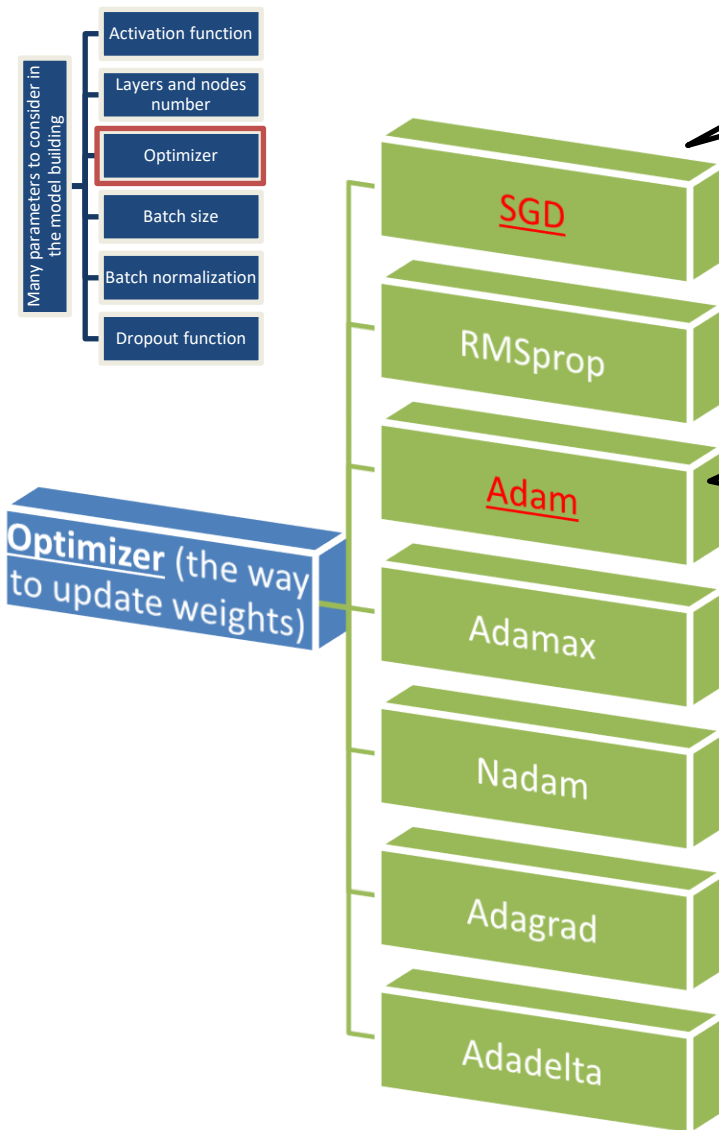
Improving the MLP with Keras – TensorFlow core



Stochastic Gradient Descent

$$w_{ij} = w_{ij} - \eta * \delta_j * x_i$$

Improving the MLP with Keras – TensorFlow core



Stochastic Gradient Descent

$$w_{ij} = w_{ij} - \eta * \delta_j * x_i$$

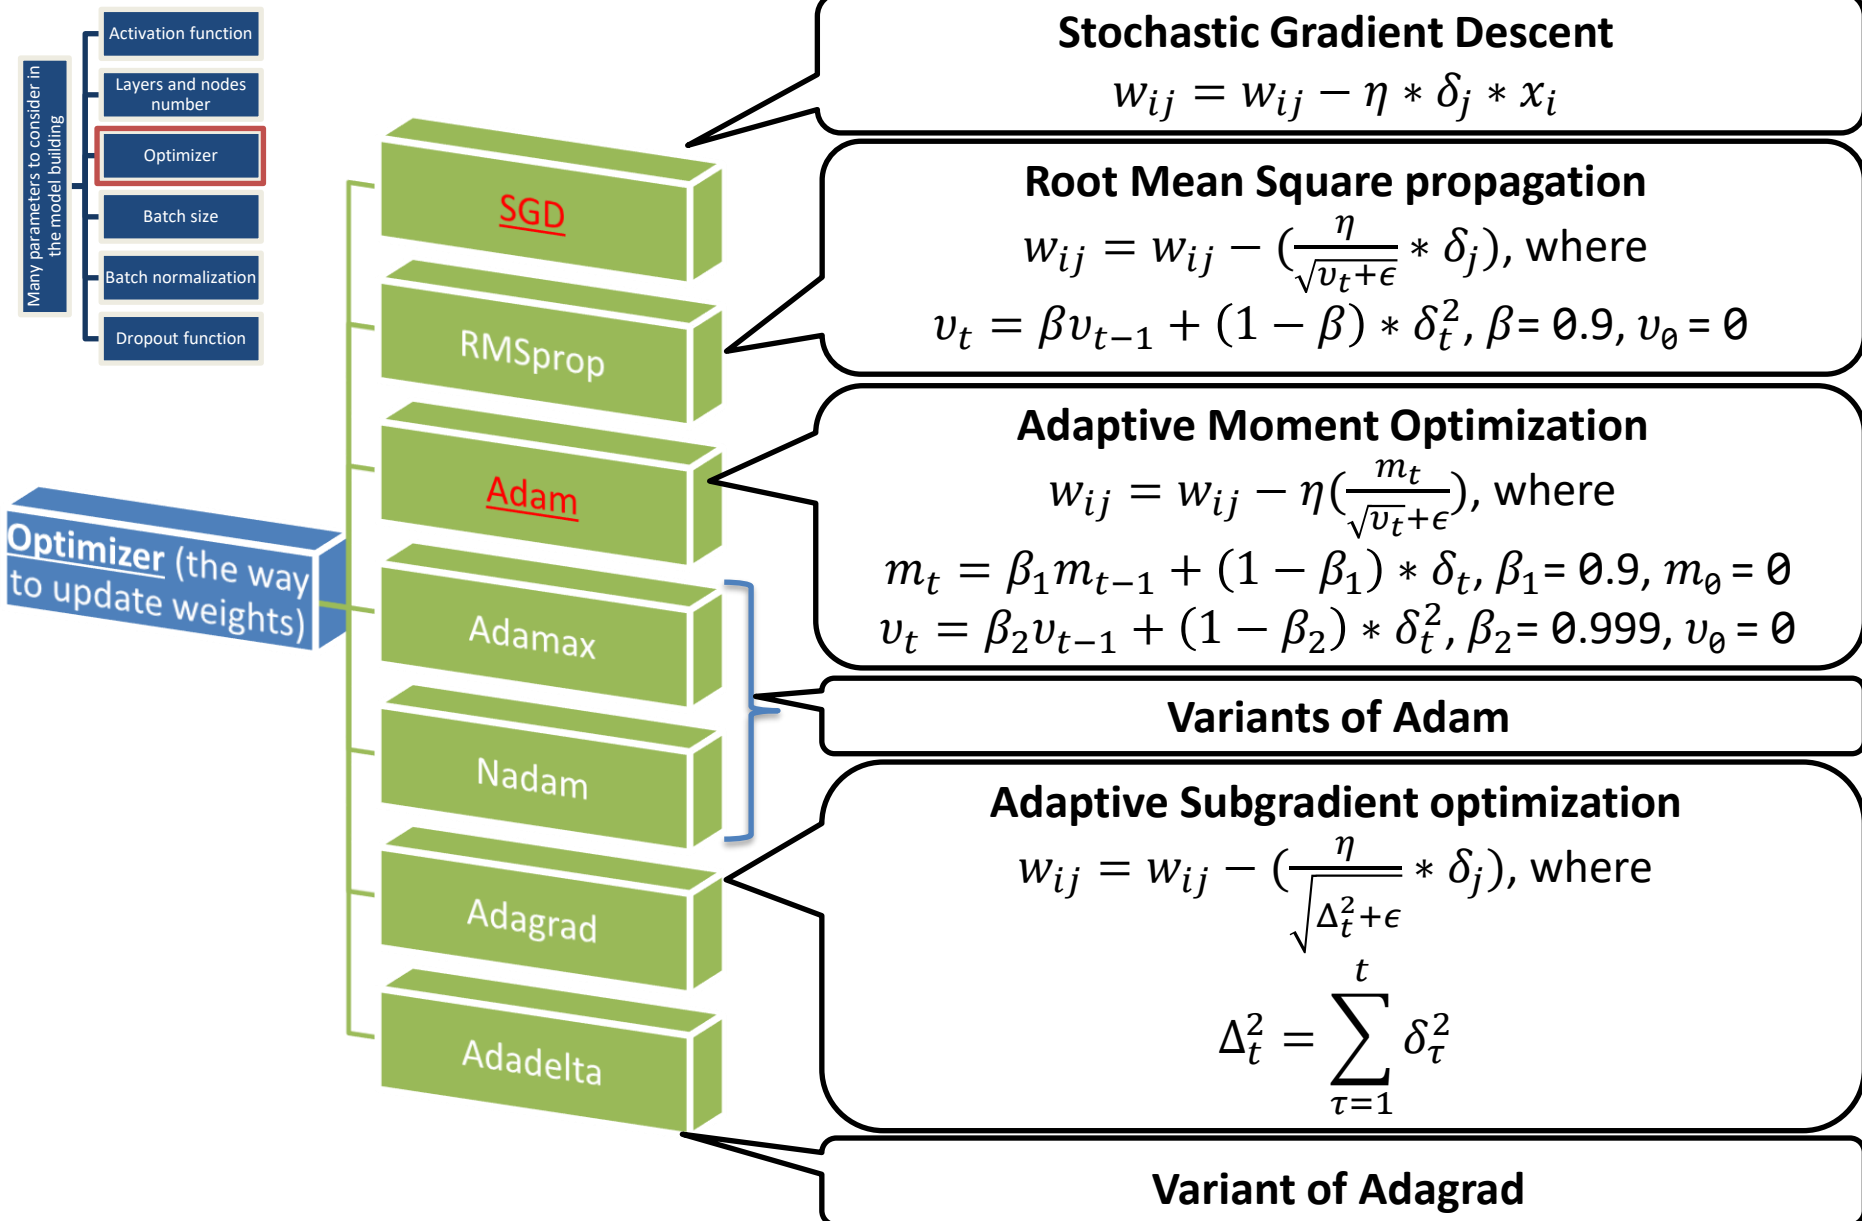
Adaptive Moment Optimization

$$w_{ij} = w_{ij} - \eta \left(\frac{m_t}{\sqrt{v_t + \epsilon}} \right), \text{ where}$$


$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) * \delta_t, \beta_1 = 0.9, m_0 = 0$$

$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) * \delta_t^2, \beta_2 = 0.999, v_0 = 0$$

Improving the MLP with Keras – TensorFlow core



Improving the MLP with Keras – TensorFlow core



Batch size

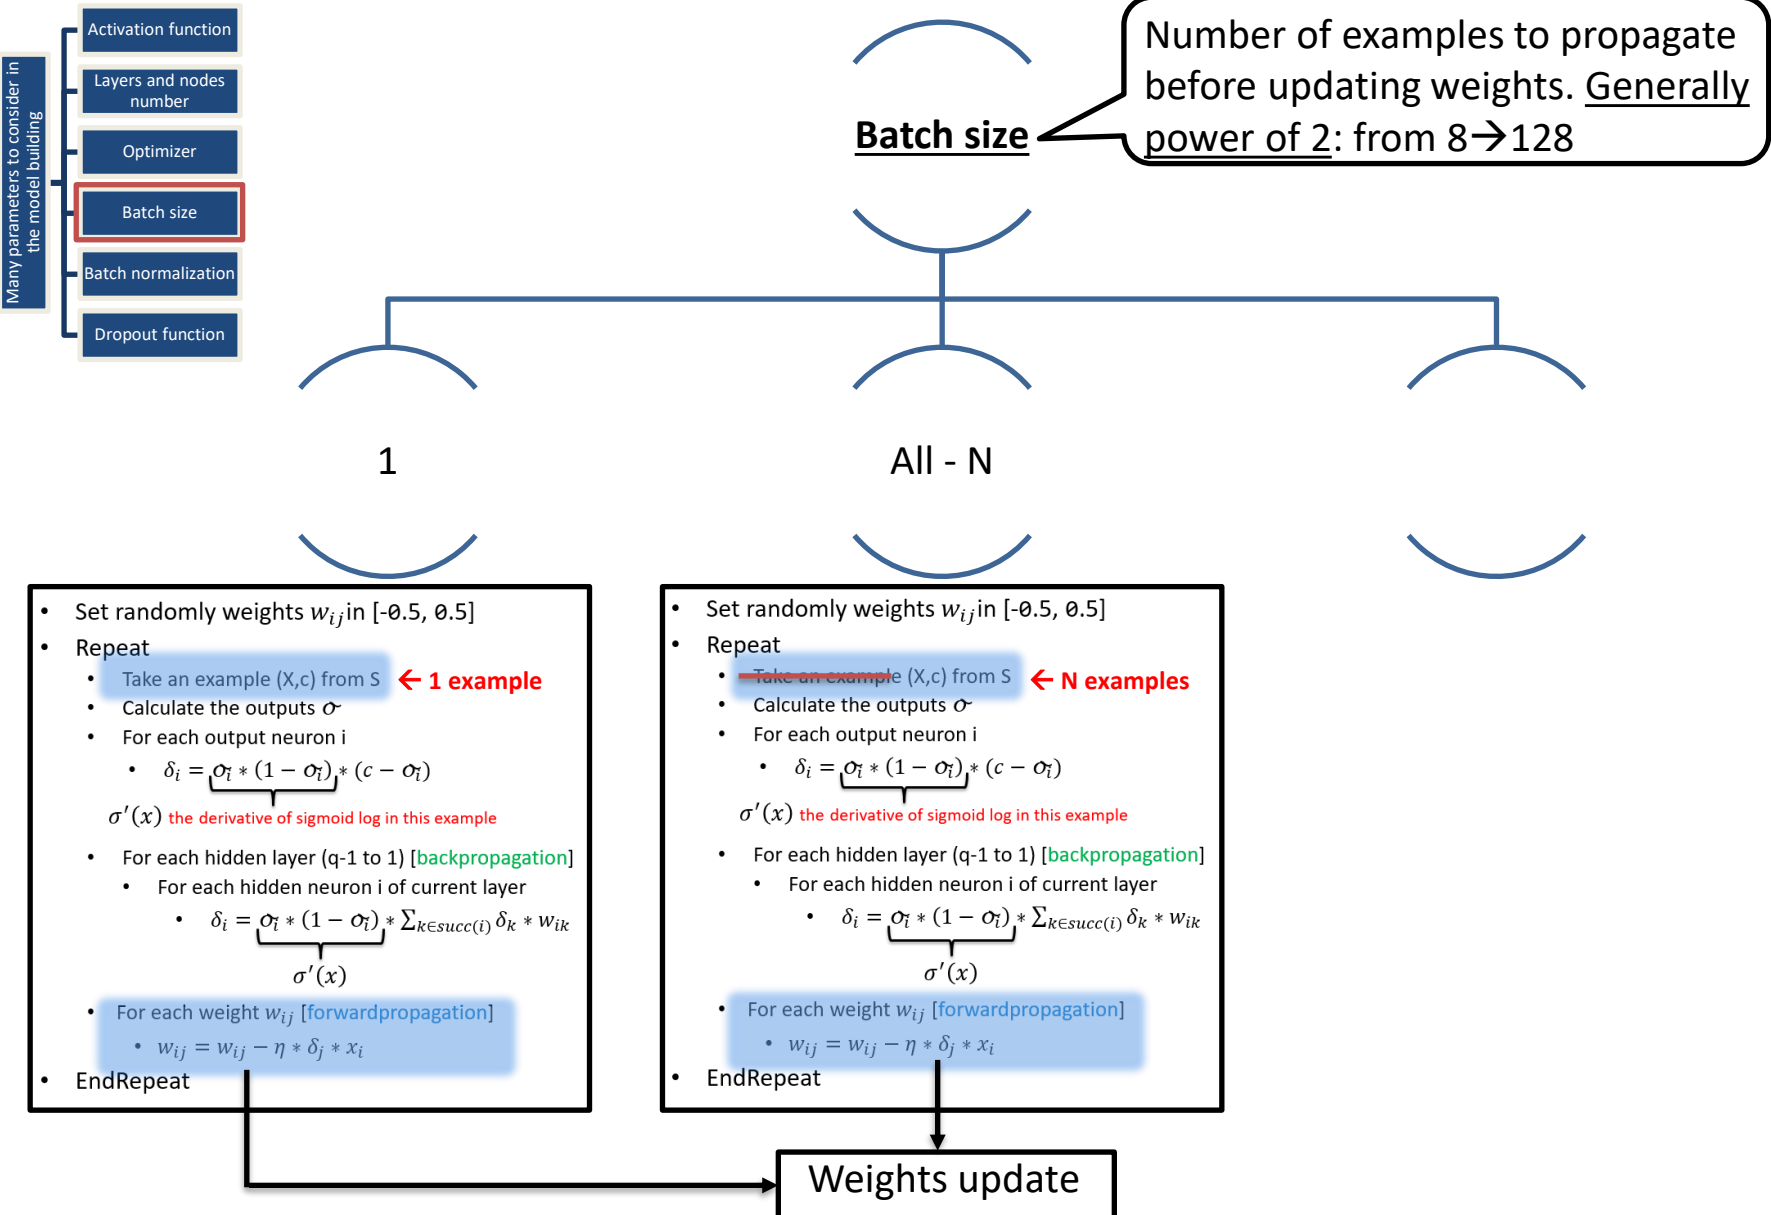
Number of examples to propagate before updating weights. Generally power of 2: from 8→128

1

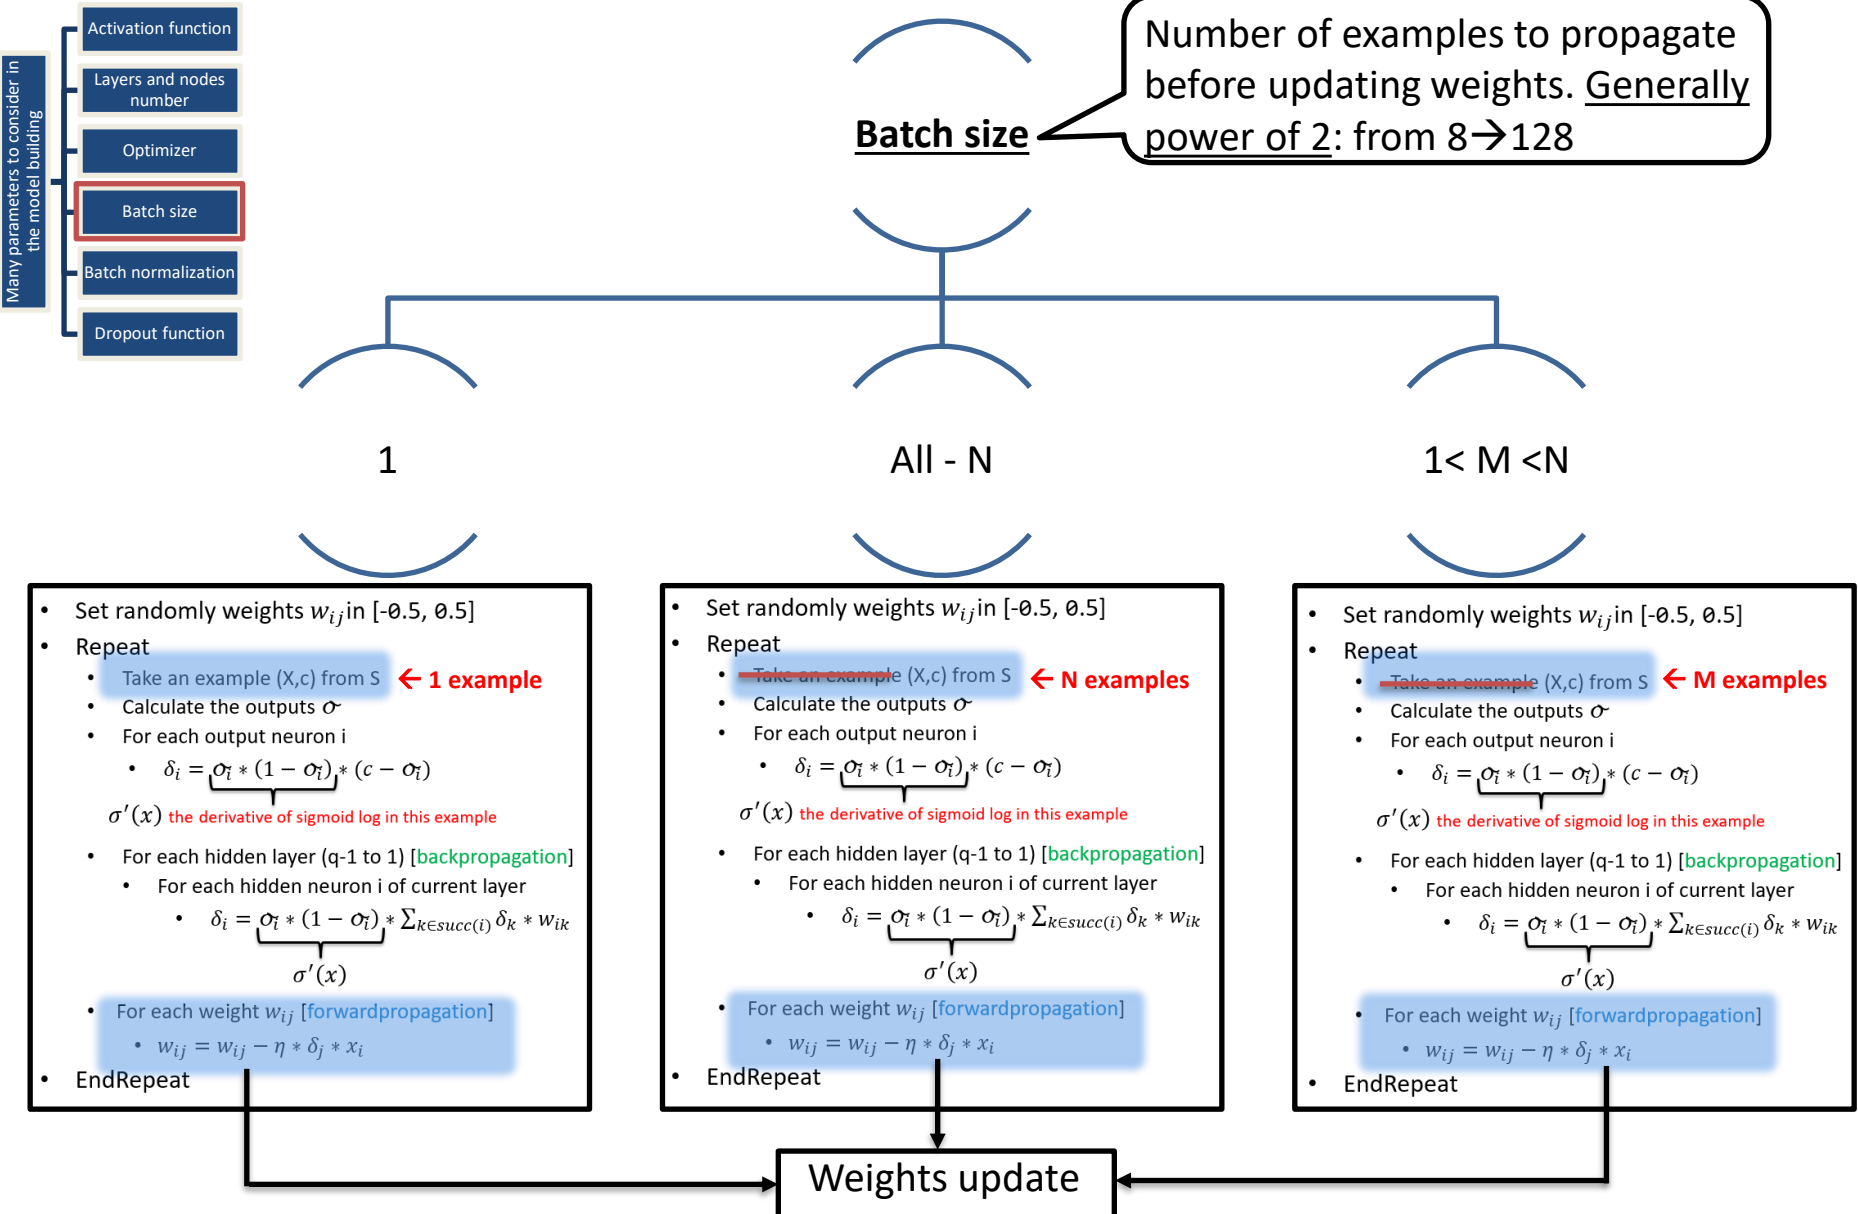
- Set randomly weights w_{ij} in $[-0.5, 0.5]$
- Repeat
 - Take an example (X, c) from S ← 1 example
 - Calculate the outputs σ
 - For each output neuron i
 - $\delta_i = \underbrace{\sigma_i * (1 - \sigma_i)}_{\sigma'(x)} * (c - \sigma_i)$
 - $\sigma'(x)$ the derivative of sigmoid log in this example
 - For each hidden layer $(q-1 \text{ to } 1)$ [backpropagation]
 - For each hidden neuron i of current layer
 - $\delta_i = \underbrace{\sigma_i * (1 - \sigma_i)}_{\sigma'(x)} * \sum_{k \in \text{succ}(i)} \delta_k * w_{ik}$
 - For each weight w_{ij} [forwardpropagation]
 - $w_{ij} = w_{ij} - \eta * \delta_j * x_i$
- EndRepeat

Weights update

Improving the MLP with Keras – TensorFlow core



Improving the MLP with Keras – TensorFlow core



Improving the MLP with Keras – TensorFlow core





Figure 2: Convergence trajectories of training and testing accuracy for SB and LB methods

Table 2: Performance of small-batch (SB) and large-batch (LB) variants of ADAM on the 6 networks listed in Table 1



Network Name	Training Accuracy		Testing Accuracy	
	SB	LB	SB	LB
F_1	99.66% \pm 0.05%	99.92% \pm 0.01%	98.03% \pm 0.07%	97.81% \pm 0.07%
F_2	99.99% \pm 0.03%	98.35% \pm 2.08%	64.02% \pm 0.2%	59.45% \pm 1.05%
C_1	99.89% \pm 0.02%	99.66% \pm 0.2%	80.04% \pm 0.12%	77.26% \pm 0.42%
C_2	99.99% \pm 0.04%	99.99 \pm 0.01%	89.24% \pm 0.12%	87.26% \pm 0.07%
C_3	99.56% \pm 0.44%	99.88% \pm 0.30%	49.58% \pm 0.39%	46.45% \pm 0.43%
C_4	99.10% \pm 1.23%	99.57% \pm 1.84%	63.08% \pm 0.5%	57.81% \pm 0.17%



<https://stats.stackexchange.com/questions/164876/what-is-the-trade-off-between-batch-size-and-number-of-iterations-to-train-a-neu>

Improving the MLP with Keras – TensorFlow core

- **Batch normalization**
 - Comes after each dense hidden layer



<https://www.youtube.com/watch?v=dXB-KQYkzNU>

Input: Values of x over a mini-batch: $\mathcal{B} = \{x_{1...m}\}$;

Parameters to be learned: γ, β

Output: $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$

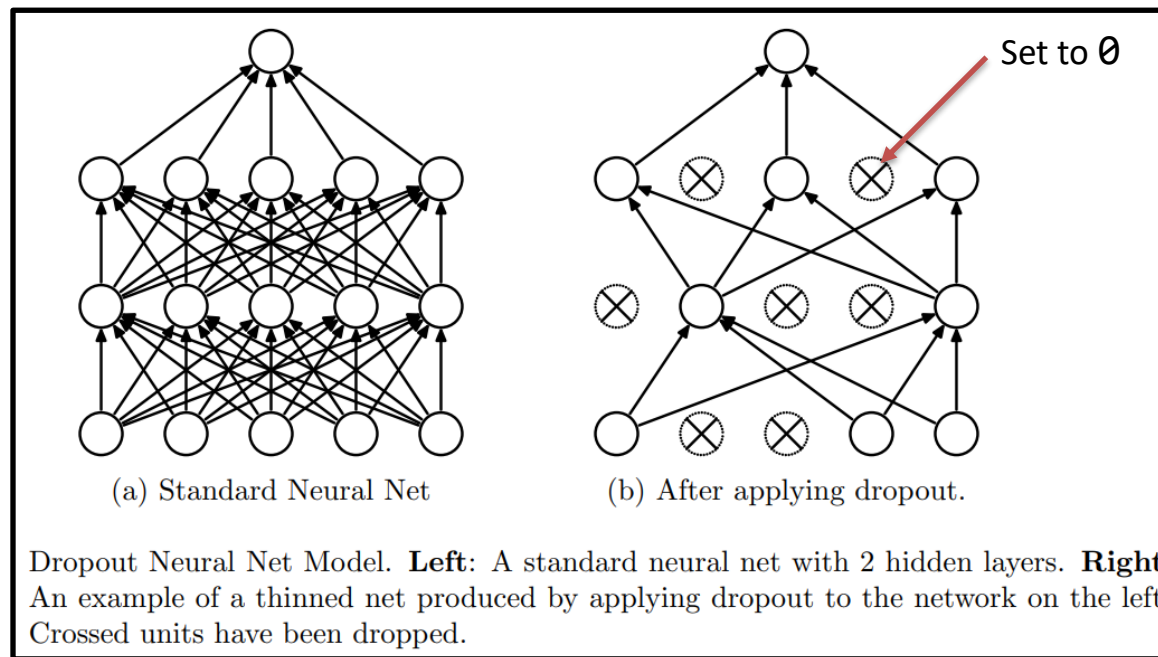

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$

Algorithm 1: Batch Normalizing Transform, applied to activation x over a mini-batch.

Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift - <https://arxiv.org/pdf/1502.03167v3.pdf>

Improving the MLP with Keras – TensorFlow core

- **Dropout function**
 - Comes after batch normalization

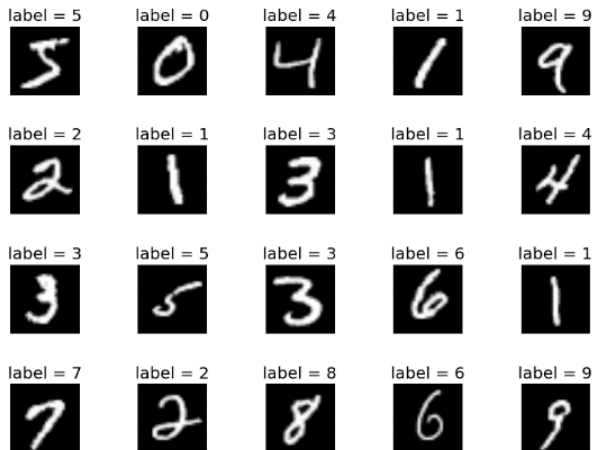


Decreasing probability p through layers. From $0.8 - 0.95 \rightarrow 0.1$

Dropout: A Simple Way to Prevent Neural Networks from Overfitting -
<http://www.jmlr.org/papers/volume15/srivastava14a/srivastava14a.pdf>

Lab session Keras – TensorFlow core

- Build an MLP to solve the hand written digits problem
 - Use MNIST dataset



MNIST database of handwritten digits

Dataset of 60,000 28x28 grayscale images of the 10 digits, along with a test set of 10,000 images.

Usage:

```
(x_train, y_train), (x_test, y_test) = mnist.load_data()
```

<https://keras.io/datasets/#mnist-database-of-handwritten-digits>

Lab session Keras – TensorFlow core

- Build an MLP to solve the hand written digits problem
 - Use MNIST dataset

```
import tensorflow as tf
import numpy as np
from tensorflow.keras.utils import to_categorical

# Load data
mnist = tf.keras.datasets.mnist

# the data, split between train and validation sets
(x_train, y_train), (x_validation, y_validation) = mnist.load_data()

# convert images into one dimension from 28x28 pixels
x_train = x_train.reshape(60000, 784)
x_validation = x_validation.reshape(10000, 784)
x_train = x_train.astype('float32')
x_validation = x_validation.astype('float32')

# normalize into [0,1]
x_train /= 255
x_validation /= 255
print('train samples', x_train.shape)
print('test samples', x_validation.shape)

print('train label samples', y_train.shape)
print('test label samples', y_validation.shape)
```

Lab session Keras – TensorFlow core

- Save your best model during training and keep history of training metrics

```
from tensorflow.keras.callbacks import ModelCheckpoint

checkpointer = ModelCheckpoint(filepath='model.hdf5', monitor='val_loss', verbose=1, save_best_only=True)

# compile the model with adam
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# train model with fit for 10 epochs
history = model.fit(x_train, y_train, validation_data = (x_validation, y_validation), epochs=10, callbacks = [checkpointer])
```

Lab session Keras – TensorFlow core

- Using the categorical cross entropy loss 'categorical_crossentropy'
https://keras.io/api/losses/probabilistic_losses/#categorical_crossentropy-class
 - Transform your label data in one_hot representation

```
# if we want to use categorical_crossentropy loss instead of sparse we have to convert the "y" labels  
from tensorflow.keras.utils import to_categorical  
  
y_train = tf.keras.utils.to_categorical(y_train, num_classes=10)  
y_validation = tf.keras.utils.to_categorical(y_test, num_classes=10)
```

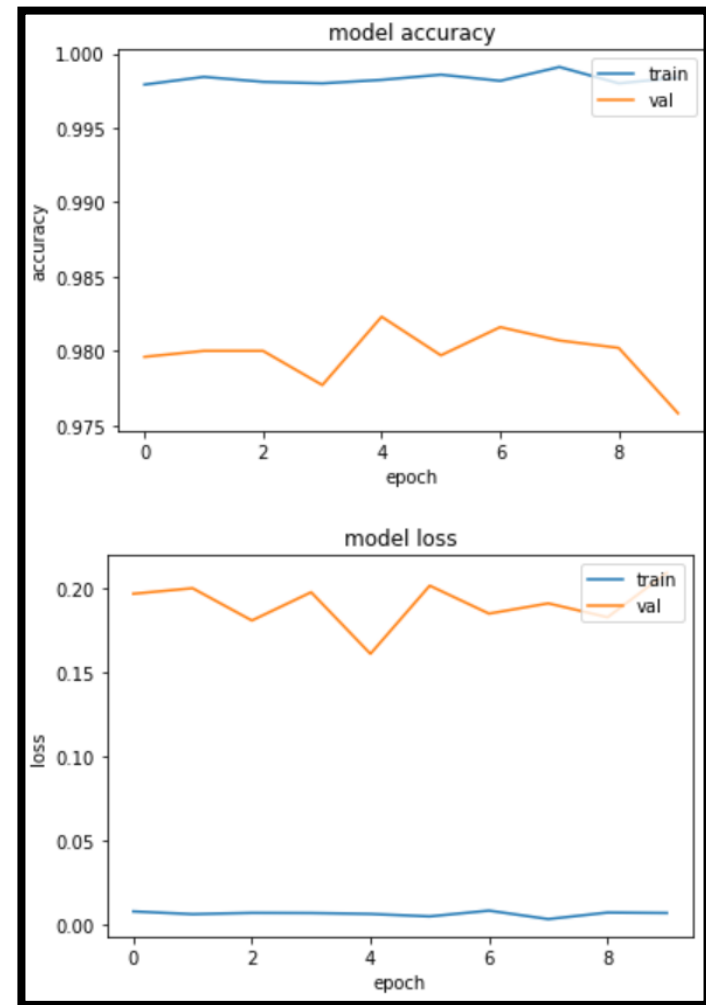
Lab session Keras – TensorFlow core

- Show training and validation curves

```
from matplotlib import pyplot as plt

plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'val'], loc='upper right')
plt.show()

plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'val'], loc='upper right')
plt.show()
```




Lab session Keras – TensorFlow core

- Load model and test it on an image

```
from tensorflow.keras.models import load_model  
  
model = load_model('model.hdf5')
```

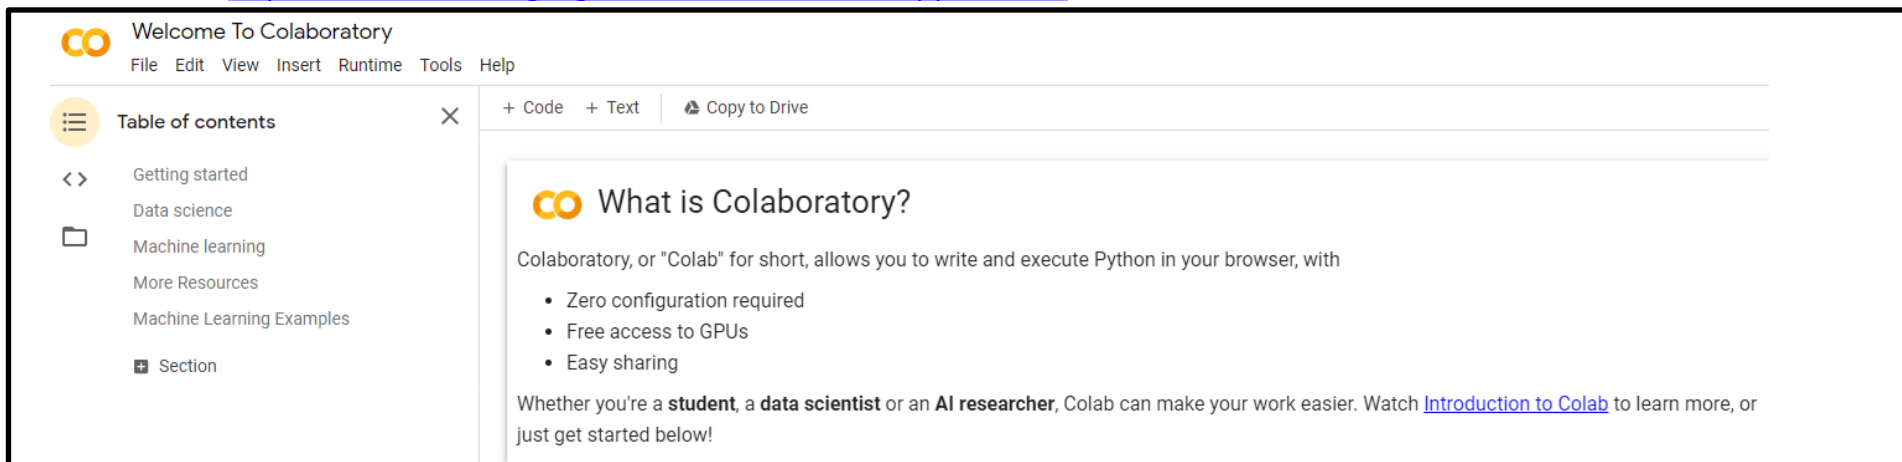
```
import cv2  
from matplotlib import pyplot as plt  
  
image = x_validation[6].reshape(28,28)  
plt.imshow(cv2.cvtColor(image, cv2.COLOR_BGR2RGB))  
plt.axis('off')  
  
datapoint = x_validation[6].reshape(1,784)  
predict_prob = model.predict(datapoint)  
print('I am confident around {:.2f}% that this image corresponds to digit {}'.format(np.amax(predict_prob)*100,  
                                                                                      np.argmax(predict_prob)))
```

```
I am confident around 100.00% that this image corresponds to digit 4
```



Lab session Keras – TensorFlow core

- Give the performance of your MLP model (same architecture) by
 - Using the categorical cross entropy loss 'categorical_crossentropy'
https://keras.io/api/losses/probabilistic_losses/#categorical_crossentropy-class
 - Testing three batch sizes: 16, 32, 64
 - Testing two optimizers: SGD, ADAM
 - with/without batch normalization and dropout
 - Show training and validation accuracy/loss curves for each scenario
- If your machine is not equipped with a GPU
 - Use the online Google Colaboratory platform. In the “Edit” option activate a GPU
<https://colab.research.google.com/notebooks/intro.ipynb?hl=en>



The screenshot displays the Google Colaboratory web interface. At the top, there's a 'Welcome To Colaboratory' header with a navigation menu (File, Edit, View, Insert, Runtime, Tools, Help). Below this, a sidebar on the left contains a 'Table of contents' with links to 'Getting started', 'Data science', 'Machine learning', 'More Resources', and 'Machine Learning Examples'. The main content area shows a document titled 'What is Colaboratory?' with the Colab logo. The text explains that Colab allows writing and executing Python in the browser, highlighting features like zero configuration, free GPU access, and easy sharing. It also mentions that it's useful for students, data scientists, and AI researchers, and provides a link to the 'Introduction to Colab'.